

Aim -

To study and implement Security as a Service on AWS/Azure.

Theory -

A business model called SECaaS, or Security as a Service, offers security to IT companies on a subscription basis. A superior security platform is provided by the outsourced approach, which lowers the total cost of ownership than the business could supply on its own. With the use of cloud computing, security for the company is maintained by an outside party. For the necessary computational and storage resources to run their websites and apps, many enterprises rely on security services. SECaaS is impressed by the “Security as a Service (SaaS)” model as applied to implement security kind services and doesn’t need on-premises hardware, avoiding substantial capital outlays. These security services typically embody authentication, antivirus, anti-malware/spyware, intrusion detection, penetration testing, and security event management, among others.

The former method of doing things involved paying direct pricing for hardware as well as ongoing fees for licenses to allow for the usage of that security code, which made it much more expensive. Instead, security as a service makes it simple and rational to use similar technologies.

Security can be availed by the following alliances:

- Encryption: makes the data unreadable until it has been authentically decoded, or encrypted.
- Network security: Network access management protocols are used to secure and keep an eye on network services.
- Email security: Protects against email frauds, spam, phishing, malware etc.
- Identification: Users can access with a valid login ID and legal permission, else forbids if it is not authenticated.
- Data loss prevention: Tools are built to monitor and secure data to protect from data loss.

Procedure-

Login to your AWS Management Console.

Create a new AWS EC2 instance with Ubuntu OS.

Screenshot of the AWS Cloud9 interface showing the selection of an Amazon Machine Image (AMI). The selected AMI is "Ubuntu Server 22.04 LTS (HVM), SSD Volume Type" (ami-080e1f13689e07408). A modal window titled "Summary" provides details about the instance, including its type (t2.micro), security group (New security group), and storage (1 volume(s) - 8 GiB). A notification indicates a "Free tier: In your first year" discount.

Screenshot of the AWS EC2 Instances page showing the successful launch of an instance. The instance ID is i-040ca33b7add83dbe. The "Next Steps" section offers links to various EC2 management tasks like creating alerts, connecting to instances, and configuring RDS databases.

Screenshot of the AWS EC2 Dashboard showing the list of instances. One instance, named "mikil" with ID i-040ca33b7add83dbe, is listed as "Running" on a t2.micro instance type. The "Select an instance" dropdown menu is open, indicating the user is about to interact with this specific instance.

Connect with the EC2 instance using the AWS built in console.

The screenshot shows the AWS EC2 Instance Connect interface. At the top, there's a navigation bar with tabs for Cloud9, EC2, Elastic Beanstalk, AWS Amplify, Lambda, S3, and API Gateway. Below the navigation bar, there are four tabs: EC2 Instance Connect (selected), Session Manager, SSH client, and EC2 serial console. The main area displays the following information:

- Instance ID:** i-040ca33b7add83dbe (mikil)
- Connection Type:** A radio button is selected for "Connect using EC2 Instance Connect" (with the note "Connect using the EC2 Instance Connect browser-based client, with a public IPv4 address"). An alternative option "Connect using EC2 Instance Connect Endpoint" (with the note "Connect using the EC2 Instance Connect browser-based client, with a private IPv4 address and a VPC endpoint") is also shown.
- Public IP address:** 54.161.70.77
- Username:** ubuntu (with a note: "Enter the username defined in the AMI used to launch the instance. If you didn't define a custom username, use the default username, ubuntu.")
- Note:** "Note: In most cases, the default username, ubuntu, is correct. However, read your AMI usage instructions to check if the AMI owner has changed the default AMI username."

At the bottom right of this panel are "Cancel" and "Connect" buttons. Below this panel, there's a terminal window showing the following output:

```
aws [Alt+S] N. Virginia M. Mikil Lalwani
Cloud9 EC2 Elastic Beanstalk AWS Amplify Lambda S3 API Gateway
Swap usage: 0%
Expanded Security Maintenance for Applications is not enabled.
0 updates can be applied immediately.
Enable ESM Apps to receive additional future security updates.
See https://ubuntu.com/esm or run: sudo pro status

The list of available updates is more than a week old.
To check for new updates run: sudo apt update

The programs included with the Ubuntu system are free software;
the exact distribution terms for each program are described in the
individual files in /usr/share/doc/*copyright.

Ubuntu comes with ABSOLUTELY NO WARRANTY, to the extent permitted by
applicable law.

To run a command as administrator (user "root"), use "sudo <command>".
See "man sudo_root" for details.

ubuntu@ip-172-31-84-45:~$
```

The terminal window shows the user's session details: i-040ca33b7add83dbe (mikil) and Public IPs: 54.161.70.77 Private IPs: 172.31.84.45.

General Hardening

In this first step, you will implement some initial hardening configurations to improve the overall security of your SSH server.

The exact hardening configuration that is most suitable for your own server depends heavily on your own threat model and risk threshold. However, the configuration you'll use in this step is a general secure configuration that will suit the majority of servers. You will edit the main OpenSSH configuration file in /etc/ssh/sshd_config to set the majority of the hardening options in this tutorial. Before continuing it is a good idea to create a backup of your existing configuration file so that you can restore it in the unlikely event that something goes wrong.

Create a backup of the file using the following cp command:

```
sudo cp /etc/ssh/sshd_config /etc/ssh/sshd_config.bak
```

This will save a backup copy of the file to /etc/ssh/sshd_config.bak.

```
ubuntu@ip-172-31-84-45:~$ sudo cp /etc/ssh/sshd_config /etc/ssh/sshd_config.bak
ubuntu@ip-172-31-84-45:~$
```

Next, review the current default OpenSSH configuration options that correspond to the settings in /etc/ssh/sshd_config. To do this, run the following command:

```
sudo sshd -T
```

This will run OpenSSH server in extended test mode, which will validate the full configuration file and print out the effective configuration values.



```
aws Services Search [Alt+S] N. Virginia Mlik Lalwani
Cloud9 EC2 Elastic Beanstalk AWS Amplify Lambda S3 API Gateway
9,ecdsa-sha2-nistp256,ecdsa-sha2-nistp384,ecdsa-sha2-nistp521,sk-ssh-ed25519@openssh.com,sk-ecdsa-sha2-nistp256@openssh.com,rsa-sha2-512,rsa-sha2-256
pubkeyacceptedalgorithms ssh-ed25519-cert-v01@openssh.com,ecdsa-sha2-nistp256-cert-v01@openssh.com,ecdsa-sha2-nistp384-cert-v01@openssh.com,ecdsa-sha2-nistp521-cert-v01@openssh.com,sk-ssh-ed25519-cert-v01@openssh.com,sk-ecdsa-sha2-nistp256-cert-v01@openssh.com,rsa-sha2-512-cert-v01@openssh.com,rsa-sha2-256-cert-v01@openssh.com,ssh-ed25519,ecdsa-sha2-nistp256,ecdsa-sha2-nistp384,ecdsa-sha2-nistp521,sk-ssh-ed25519@openssh.com,sk-ecdsa-sha2-nistp256@openssh.com,rsa-sha2-512,rsa-sha2-256
LogLevel INFO
syslogfacility AUTH
authorizedkeysfile .ssh/authorized_keys .ssh/authorized_keys2
hostkey /etc/ssh/ssh_host_rsa_key
hostkey /etc/ssh/ssh_host_ecdsa_key
hostkey /etc/ssh/ssh_host_ed25519_key
acceptenv LANG
acceptenv LC_ALL
authenticationmethods any
subsystem sftp /usr/lib/openssh/sftp-server
maxstartups 10:30:100
personcmaxstartups none
personcnetblocksize 32:128
permittunnel no
ipgosh lowdelay throughput
rekeylimit 0 0
permitopen any
permitlisten any
permituserenvironment no
pubkeyauthoptions none
ubuntu@ip-172-31-84-45:~$
```

You can now open the configuration file using nano or your favorite text editor to begin implementing the initial hardening measures:

```
sudo nano /etc/ssh/sshd_config
```

The first hardening option is to disable logging in via SSH as the root user. Set the PermitRootLogin option to no by uncommenting or editing the line in sshd_config:

```
PermitRootLogin no
```

This option will prevent a potential attacker from logging into your server directly as root. It also encourages good operational security practices on your part, such as operating as a non-privileged user and using sudo to escalate privileges only when absolutely needed.

```
#LoginGraceTime 2m
PermitRootLogin no
#StrictModes yes
#MaxAuthTries 6
#MaxSessions 10
```

Next, you can limit the maximum number of authentication attempts for a particular login session by configuring the MaxAuthTries option in sshd_config

```
MaxAuthTries 3
```

A standard value of 3 is acceptable for most setups, but you may wish to set this higher or lower depending on your own risk threshold.

```
# Authentication:  
  
#LoginGraceTime 2m  
PermitRootLogin no  
#StrictModes yes  
MaxAuthTries 3  
File to insert [from ./]:  
^G Help M-F New Buffer  
^C Cancel M-N No Conversion
```

If required, you can also set a reduced login grace period, which is the amount of time a user has to complete authentication after initially connecting to your SSH server in `sshd_config`

`LoginGraceTime 20m`

The configuration file specifies this value in seconds.

```
#LogLevel INFO
```

```
# Authentication:
```

```
LoginGraceTime 20m  
PermitRootLogin no  
#StrictModes yes  
MaxAuthTries 3  
#MaxSessions 10
```

```
^G Help ^O Write Out  
^X Exit ^R Read File
```

Setting this to a lower value helps to prevent certain denial-of-service attacks where multiple authentication sessions are kept open for a prolonged period of time.

If you have configured SSH keys for authentication, rather than using passwords, disable SSH password authentication to prevent leaked user passwords from allowing an attacker to log in `sshd_config`

`PasswordAuthentication no`

```
# Don't read the user's ~/.rhosts and ~/.shosts files  
#IgnoreRhosts yes  
  
# To disable tunneled clear text passwords, change to no here!  
PasswordAuthentication no  
#PermitEmptyPasswords no  
  
# Change to yes to enable challenge-response passwords (beware issues with
```

```
^G Help ^O Write Out ^W Where Is ^K Cut ^T Execute ^C  
^X Exit ^R Read File ^M Replace ^U Paste ^J Justify ^/
```

As a further hardening measure related to passwords, you may also wish to disable authentication with empty passwords. This will prevent logins if a user's password is set to a blank or empty value in `sshd_config`

`PermitEmptyPasswords no`

In the majority of use cases, SSH will be configured with public key authentication as the only in-use authentication method. However, OpenSSH server also supports many other authentication methods, some of which are enabled by default. If these are not

required, you can disable them to further reduce the attack surface of your SSH server in `sshd_config`

`ChallengeResponseAuthentication no`

`KerberosAuthentication no`

`GSSAPIAuthentication no`

```
# To disable tunneled clear text passwords, change to no here!
PasswordAuthentication no
#PermitEmptyPasswords no

# Change to yes to enable challenge-response passwords (beware issues with
# some PAM modules and threads)
KbdInteractiveAuthentication no

# Kerberos options
KerberosAuthentication no
#KerberosOrLocalPasswd yes
#KerberosTicketCleanup yes
#KerberosGetAFSToken no

# GSSAPI options
GSSAPIAuthentication no
#GSSAPICleanupCredentials yes
#GSSAPIStrictAcceptorCheck yes
```

^G Help **^O Write Out** **^W Where Is** **^K Cut** **^T Execute**

X11 forwarding allows for the display of remote graphical applications over an SSH connection, but this is rarely used in practice. Disable it if you are not running a graphical environment on your server in `sshd_config`

`X11Forwarding no`

```
#AllowTcpForwarding yes
#GatewayPorts no
X11Forwarding no
#X11DisplayOffset 10
```

^G Help **^O Write Out** **^W Where Is** **^K Cut** **^T Execute**

OpenSSH server allows connecting clients to pass custom environment variables. For example, a client can attempt to set its own \$PATH or to configure terminal settings. However, like X11 forwarding, these are not commonly used, so you can disable the option in most cases in `sshd_config`

`PermitUserEnvironment no`

```
#TCPKeepAlive yes
PermitUserEnvironment no
#Compression delayed
```

^G Help **^O Write Out** **^W 1**

Implementing an IP Address Allowlist

You can use IP address allow lists to limit the users who are authorized to log in to your server on a per-IP address basis. In this step, you will configure an IP allowlist for your

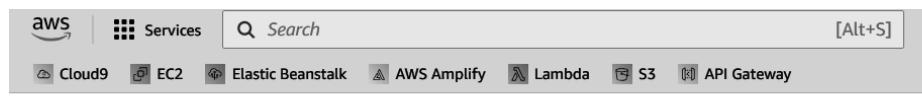
OpenSSH server.

In many cases, you will only be logging on to your server from a small number of known, trusted IP addresses. For example, your home internet connection, a corporate VPN appliance, or a static jump box or bastion host in a data center.

By implementing an IP address allowlist, you can ensure that people will only be able to log in from one of the pre-approved IP addresses, greatly reducing the risk of a breach in the event that your private keys and/or passwords are leaked.

You can identify the IP address that you're currently connecting to your server with by using the w command:

W



```
aws | Services | Search [Alt+S]
Cloud9 EC2 Elastic Beanstalk AWS Amplify Lambda S3 API Gateway
ubuntu@ip-172-31-84-45:~$ w
07:05:57 up 2:38, 1 user, load average: 0.00, 0.00, 0.00
USER     TTY      FROM          LOGIN@    IDLE   JCPU   PCPU WHAT
ubuntu   pts/0    18.206.107.28  06:21     2.00s  0.02s  0.00s w
ubuntu@ip-172-31-84-45:~$ █
```

This will output something similar to the following:

Locate your user account in the list and take a note of the connecting IP address. Here we use the example IP of 13.233.177.5

In order to begin implementing your IP address allowlist, open the OpenSSH server configuration file in nano or your preferred text editor:

`sudo nano /etc/ssh/sshd_config`

You can implement IP address allowlists using the `AllowUsers` configuration directive, which restricts user authentications based on username and/or IP address.

Your own system setup and requirements will determine which specific configuration is the most appropriate. The following examples will help you to identify the most suitable one:

Restrict all users to a specific IP address:

`AllowUsers *@18.206.107.28`

```
# Example of overriding settings on a per-user basis
#Match User anoncvs
#       X11Forwarding no
#       AllowTcpForwarding no
#       PermitTTY no
#       ForceCommand cvs server
AllowUsers *@18.206.107.28█
```

Restrict all users to a specific IP address range using Classless Inter-Domain Routing (CIDR) notation:

`AllowUsers *@18.206.107.28/24`

```
#      X11Forwarding no
#      AllowTcpForwarding no
#      PermitTTY no
#      ForceCommand cvs server
AllowUsers *@18.206.107.28/24
```

Restrict all users to a specific IP address range (using wildcards):

AllowUsers *@18.206.107.*

```
#      X11Forwarding no
#      AllowTcpForwarding no
#      PermitTTY no
#      ForceCommand cvs server
AllowUsers *@18.206.107.28.*
```

^G Help **^C Write Out** **^W Where I**

Restrict all users to multiple specific IP addresses and ranges:

AllowUsers *@18.206.107.28 *@18.206.107.29 *@192.0.2.0/24 *@172.16.*.1

```
#      X11Forwarding no
#      AllowTcpForwarding no
#      PermitTTY no
#      ForceCommand cvs server
AllowUsers *@18.206.107.28 *@18.206.107.29 *@192.0.2.0/24 *@172.16.*.1
```

^G Help **^C Write Out** **^W Where Is** **^K Cut** **^T Execute**

Disallow all users except for named users from specific IP addresses:

AllowUsers sammy@18.206.107.28 alex@18.206.107.29

```
#      X11Forwarding no
#      AllowTcpForwarding no
#      PermitTTY no
#      ForceCommand cvs server
AllowUsers sammy@18.206.107.28 alex@18.206.107.29
```

Restrict a specific user to a specific IP address, while continuing to allow all other users to log in without restrictions:

Match User ashley

AllowUsers sammy@18.206.107.28

```
#      X11Forwarding no
#      AllowTcpForwarding no
#      PermitTTY no
#      ForceCommand cvs server
AllowUsers sammy@18.206.107.28
```

Save and close the file, and then proceed to test your configuration syntax:

```
sudo sshd -t
ubuntu@ip-172-31-84-45:~$ sudo nano /etc/ssh/sshd_config
ubuntu@ip-172-31-84-45:~$ sudo sshd -t
ubuntu@ip-172-31-84-45:~$
```

If no errors are reported, you can reload OpenSSH server to apply your configuration:

```
sudo systemctl reload sshd.service
```

```
ubuntu@ip-172-31-84-45:~$ sudo nano /etc/ssh/sshd_config
ubuntu@ip-172-31-84-45:~$ sudo sshd -t
ubuntu@ip-172-31-84-45:~$ sudo systemctl reload sshd.service
ubuntu@ip-172-31-84-45:~$
```

Restricting the Shell of a User

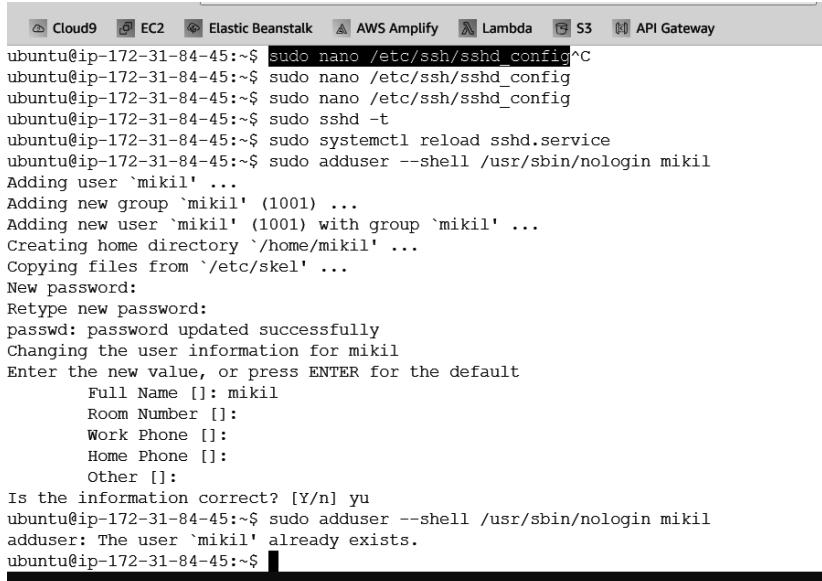
In this step, you'll explore the various options for restricting the shell of an SSH user. In addition to providing remote shell access, SSH is also great for transferring files and other data, for example, via SFTP. However, you may not always want to grant full shell access to users when they only need to be able to carry out file transfers.

There are multiple configurations within OpenSSH server that you can use to restrict the shell environment of particular users. For instance, in this tutorial, we will use these to create SFTP-only users.

Firstly, you can use the /usr/sbin/nologin shell to disable interactive logins for certain user accounts, while still allowing non-interactive sessions to function, like file transfers, tunneling, and so on.

To create a new user with the nologin shell, use the following command:

```
sudo adduser --shell /usr/sbin/nologin mikil
```



A screenshot of the AWS Cloud9 terminal interface. The top navigation bar includes links for Cloud9, EC2, Elastic Beanstalk, AWS Amplify, Lambda, S3, and API Gateway. The terminal window displays a series of commands run on an Ubuntu system (ip-172-31-84-45). The user creates a new user account named 'mikil' using the sudo adduser command. The process involves setting a password, entering contact information (Full Name, Room Number, Work Phone, Home Phone, Other), and confirming the user already exists. The terminal shows the user creation command and its output.

```
ubuntu@ip-172-31-84-45:~$ sudo nano /etc/ssh/sshd_config^C
ubuntu@ip-172-31-84-45:~$ sudo nano /etc/ssh/sshd_config
ubuntu@ip-172-31-84-45:~$ sudo nano /etc/ssh/sshd_config
ubuntu@ip-172-31-84-45:~$ sudo sshd -t
ubuntu@ip-172-31-84-45:~$ sudo systemctl reload sshd.service
ubuntu@ip-172-31-84-45:~$ sudo adduser --shell /usr/sbin/nologin mikil
Adding user `mikil' ...
Adding new group `mikil' (1001) ...
Adding new user `mikil' (1001) with group `mikil' ...
Creating home directory `/home/mikil' ...
Copying files from `/etc/skel' ...
New password:
Retype new password:
passwd: password updated successfully
Changing the user information for mikil
Enter the new value, or press ENTER for the default
  Full Name []: mikil
  Room Number []:
  Work Phone []:
  Home Phone []:
  Other []:
Is the information correct? [Y/n] yu
ubuntu@ip-172-31-84-45:~$ sudo adduser --shell /usr/sbin/nologin mikil
adduser: The user `mikil' already exists.
ubuntu@ip-172-31-84-45:~$
```

If you then attempt to interactively log in as one of these users, the request will be rejected:

```
sudo su mikil
```

This will output something similar to the following message:

```
adduser: The user `mikil' already exists.
ubuntu@ip-172-31-84-45:~$ sudo su mikil
This account is currently not available.
ubuntu@ip-172-31-84-45:~$
```

Despite the rejection message on interactive logins, other actions such as file transfers will still be allowed.

Advanced Hardening

In this final step, you will implement various additional hardening measures to make access to your SSH server as secure as possible.

OpenSSH servers can impose restrictions on a per-key basis. Specifically, restrictions can be applied to any public keys that are present in the .ssh/authorized_keys file. This ability is particularly useful to control access for machine-to-machine sessions, as well as providing the ability for non-sudo users to control the restrictions for their own user account.

While you can apply most of these restrictions at the system or user level using the /etc/ssh/sshd_configuration file, it can be advantageous to implement them at the key-level as well, to provide defense-in-depth and an additional failsafe in the event of accidental system-wide configuration errors.

Begin by opening your .ssh/authorized_keys file in nano or your preferred editor:
`nano ~/.ssh/authorized_keys`

Once you've opened your authorized_keys file, you will see that each line contains an SSH public key, which will most likely begin with something like ssh-rsa AAAB....

Additional configuration options can be added to the beginning of the line, and these will only apply to successful authentications against that specific public key. The following restriction options are available:

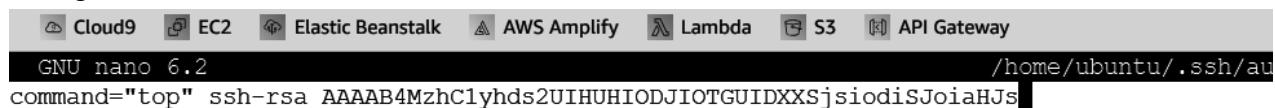
- no-agent-forwarding: Disable SSH agent forwarding.
- no-port-forwarding: Disable SSH port forwarding.
- no-pty: Disable the ability to allocate a tty (i.e. start a shell).
- no-user-rc: Prevent execution of the `~/.ssh/rc` file.
- no-X11-forwarding: Disable X11 display forwarding.

You can apply these to disable specific SSH features for specific keys. For example, to disable agent forwarding and X11 forwarding for a key, you would use the following configuration:



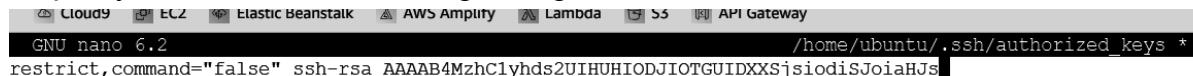
```
GNU nano 6.2 /home/ubuntu/.ssh/authorized_keys *  
no-agent-forwarding,no-X11-forwarding ssh-rsa AAAAB4MzhC1yhds2UIHUHIODJIC
```

You may also wish to consider using the command option, which is very similar to the ForceCommand option described in Step 3. This doesn't provide a direct benefit if you're already using ForceCommand, but it is good defense-in-depth to have it in place, just in the unlikely event that your main OpenSSH server configuration file is overwritten, edited, and so on. For example, to force users authenticating using a specific key to execute a specific command upon login, you can add the following configuration:



```
Cloud9 EC2 Elastic Beanstalk AWS Amplify Lambda S3 API Gateway  
GNU nano 6.2 /home/ubuntu/.ssh/authorized_keys *  
command="top" ssh-rsa AAAAB4MzhC1yhds2UIHUHIODJIOTGUIDXXSjsiodisJoi
```

Finally, to best use the per-key restrictions for the SFTP-only user that you created in Step 3, you can use the following configuration:



```
Cloud9 EC2 Elastic Beanstalk AWS Amplify Lambda S3 API Gateway  
GNU nano 6.2 /home/ubuntu/.ssh/authorized_keys *  
restrict,command="false" ssh-rsa AAAAB4MzhC1yhds2UIHUHIODJIOTGUIDXXSjsiodisJoi
```

The restrict option will disable all interactive access, and the command="false" option acts as a second line of defense in the event that the ForceCommand option or nologin shell were to fail.

Save and close the file to apply the configuration. This will take effect immediately for all new logins, so you don't need to reload OpenSSH manually.

In conclusion, you reviewed your OpenSSH server configuration and implemented various hardening measures to help secure your server. The options that you configured

have reduced the overall attack surface of your server by disabling unused features and locking down the access of specific users.

Conclusion -

Thus, we successfully studied and implemented Security as a Service on AWS/Azure.